

CVX Implementation Notes

Michael J. Meyer

December 5, 2016

1 Regularization

The quadratic approximation of the objective function $f(x)$ at iterate x_k is

$$\tilde{f}(x_k + \Delta x) = f(x_k) + \nabla f(x_k)' \Delta x + \frac{1}{2} \Delta x' H_k \Delta x,$$

where H_k is the Hessian $\nabla^2 f(x_k)$ or an approximation thereof. The search direction Δx from iterate x_k is computed by minimizing this function over the variable Δx resulting in the equation

$$H_k \Delta x = -\nabla f(x_k) := -y_k \quad (1)$$

for the search direction Δx . Clearly we can assume that $y_k = \nabla f(x_k) \neq 0$ since the search terminates at a zero gradient.

In our algorithms we can also assume that H_k is positive semidefinite, but not necessarily nonsingular. If H_k is nonsingular (i.e. positive definite), then the solution Δx is a *descent direction*:

$$\Delta x' \nabla f(x_k) = -\Delta x' H_k \Delta x < 0$$

and this is all we care about. We will solve (1) by Cholesky factorization $H_k = LL'$ with lower triangular L , by solving the triangular systems

$$Lv = -y_k, \quad L' \Delta x = v. \quad (2)$$

The Cholesky factorization fails if H_k is singular. In that case we replace H_k with $H_k(\delta) := H_k + \delta I$, for some positive constant δ . This matrix is now positive definite, in fact

$$(H_k(\delta)u, u) = u' H_k(\delta)u = u' H_k u + \delta u' I u \geq \delta \|u\|^2$$

so that (1) now yields a descent direction. Moreover it improves the conditioning of (1): if $H_k + \delta I = L(\delta)L(\delta)'$ is the Cholesky factorization of $H_k(\delta)$, then we have

$$|L(\delta)_{ii}| \geq \delta.$$

Indeed, the diagonal element $L(\delta)_{ii}$ is an eigenvalue of the triangular matrix $L(\delta)'$. Let u be a corresponding eigenvector. Then we have

$$|L(\delta)_{ii}|^2 \|u\|^2 = \|L(\delta)'u\|^2 = (L(\delta)L(\delta)'u, u) = (H_k(\delta)u, u) \geq \delta \|u\|^2$$

from which it follows that

$$|L(\delta)_{ii}| \geq \sqrt{\delta}$$

with obvious implications for the numerical stability of the triangular systems (2). Moreover this suggests that we should replace H_k with $H_k + \delta I$ not only if the Cholesky factorization fails but rather as soon as the minimal diagonal element (in absolute value) of the Cholesky factor L is below the threshold $\sqrt{\delta}$.

Trust region interpretation. The passage from the matrix H_k to the regularization $H_k + \delta I$ has an interpretation in terms of *trust regions*: the solution Δx^* of

$$H_k(\delta)\Delta x = -y_k, \quad \text{where } y_k = \nabla f(x_k),$$

is the minimizer of the quadratic function

$$\begin{aligned} \phi(\Delta x) &= f(x_k) + y'_k \Delta x + \Delta x' H_k(\delta) \Delta x \\ &= f(x_k) + y'_k \Delta x + \Delta x' H_k \Delta x + \delta \|\Delta x\|^2 \\ &= \tilde{f}(x_k + \Delta x) + \delta \|\Delta x\|^2. \end{aligned}$$

Now note that this minimizer Δx^* is automatically also the minimizer of the quadratic approximation $\tilde{f}(x_k + \Delta x)$ on the ball $B(x_k, r_k)$ with radius $r_k = \|\Delta x^*\|$. Indeed, if this ball contained a point u with $f(x_k + u) < \tilde{f}(x_k + \Delta x^*)$, then, since also $\|u\| \leq \|\Delta x^*\|$ it follows that

$$\phi(u) = \tilde{f}(x_k + u) + \delta \|u\|^2 < \tilde{f}(x_k + \Delta x^*) + \delta \|\Delta x^*\|^2 = \phi(\Delta x^*).$$

In other words: passing from H_k to $H_k + \delta I$ we compute the search direction Δx by minimizing the quadratic approximation $\tilde{f}(x_k + \Delta x)$ not globally but instead on the ball $B(x_k, r_k)$ (the region in which we trust the approximation) where the trust radius r_k is defined implicitly as $r_k = \|\Delta x^*\|$.

This indicates that the regularization $H_k \rightarrow H_k(\delta)$ is not unreasonable and in any case it solves the problem of nonsingularity of H_k for us, improves the conditioning and results in a descent direction Δx at iterate x_k .

2 Hessian

Let $f : \mathbb{R}^n \rightarrow \mathbb{R}$ be a twice continuously differentiable function. The second order Taylor expansion of f centered at x has the form

$$f(x + h) = f(x) + L(h) + \frac{1}{2}B(h, h) + R(h),$$

where L is a linear function of $h \in \mathbb{R}^n$, $B(u, v)$ is a bilinear function of $(u, v) \in \mathbb{R}^n \times \mathbb{R}^n$ and the remainder $R(h)$ satisfies $R(h) = o(\|h\|^2)$. This condition on the remainder ensures that L and B are uniquely determined as

$$L(h) = \nabla f(x)'h \quad \text{and} \quad B(u, v) = u'Hv,$$

where $H := \nabla^2 f(x) \in \text{Mat}_{n \times n}(\mathbb{R})$ is the matrix with entries

$$H_{ij} = B(e_i, e_j) = \frac{\partial^2 f}{\partial x_i \partial x_j}(x),$$

i.e. the Hessian matrix of f at x .

To compute the gradient and Hessian of a C^2 -function f we make use of the fact that the remainder condition $R(h) = o(\|h\|^2)$ in a quadratic expansion

$$f(x + h) = f(x) + \Delta' h + h' H h + R(h)$$

uniquely determines the “coefficients” Δ and H as $\Delta = \nabla f(x)$ and $H = \nabla^2 f(x)$. We only need to find such an expansion for $f(x+h)$ and check that the remainder satisfies $R(h) = o(\|h\|^2)$. This is how we will derive our formulas below.

Hessian of affine transformation. Let $f : \mathbb{R}^n \rightarrow \mathbb{R}$ be a C^2 -function and $\bar{f}(u) = f(x_0 + Fu)$, where $x_0 \in \mathbb{R}^n$ and $F : \mathbb{R}^m \rightarrow \mathbb{R}^n$ is a linear map, i.e. $F \in \text{Mat}_{n \times m}$.

We want to compute the gradient and Hessian of h at any point u from those of f . To get these do a second order Taylor expansion of f about x :

$$f(x+h) = f(x) + h^T g + \frac{1}{2} h^T H h + o(\|h\|^2),$$

where $g = \nabla f(x)$ and $H = \nabla^2 f(x)$ are uniquely determined by the fact that the residual is $o(\|h\|^2)$. Applying this to the point $x = x_0 + Fu$ this implies that

$$\begin{aligned} \bar{f}(u+h) &= f(x_0 + Fu + Fh) \\ &= f(x_0 + Fu) + (Fh)^T g + \frac{1}{2} (Fh)^T H Fh + o(\|Fh\|^2). \end{aligned}$$

Since $o(\|Fh\|^2)$ is $o(\|h\|^2)$ we conclude from this that

$$\nabla \bar{f}(u) = F^T g \quad \text{and} \quad \nabla^2 \bar{f}(u) = F^T H F,$$

or, more explicitly

$$\nabla \bar{f}(u) = F^T \nabla f(x_0 + Fu) \quad \text{and} \quad \nabla^2 \bar{f}(u) = F^T \nabla^2 f(x_0 + Fu) F. \quad (3)$$

Hessian of composition. With a similar approach we can compute the Hessian of a composition $g(f(x))$, where here $g : \mathbb{R} \rightarrow \mathbb{R}$ is a scalar function of one variable (more general g are much harder to handle and we do not need them). Indeed, set

$$y = f(x), \quad \nabla = \nabla f(x) \quad H = \nabla^2 f(x) \quad \text{and} \quad k = h^T \nabla + \frac{1}{2} h^T H h$$

and use second order Taylor approximations on f at the point x and g at the point $y = f(x)$ to obtain:

$$\begin{aligned} g(f(x+h)) &= g\left(f(x) + h^T \nabla + \frac{1}{2} h^T H h\right) \\ &= g(y+k) = g(y) + g'(y)k + \frac{1}{2} g''(y)k^2 + o(k^2) \\ &= g(y) + g'(y) \left(h^T \nabla + \frac{1}{2} h^T H h\right) + \frac{1}{2} g''(y) \left(h^T \nabla + \frac{1}{2} h^T H h\right)^2 + o(\|h\|^2). \end{aligned}$$

Here we have used that $o(k^2) = o(\|h\|^2)$. Collect terms of first and second order in h together and sticking all terms of higher order into the residual $o(\|h\|^2)$. Note that the squared term contributes no first order terms and only one second order term, this being the term

$$(h^T \nabla)^2 = (h^T \nabla)(h^T \nabla) = (h^T \nabla)(h^T \nabla)^T = h^T (\nabla \nabla^T) h.$$

We obtain

$$g(f(x+h)) = g(y) + h^T [g'(y)\nabla] + \frac{1}{2} h^T [g'(y)H + g''(y)\nabla\nabla^T] h + o(\|h\|^2).$$

and from this we can read off that

$$\nabla(g \circ f)(x) = g'(f(x))\nabla f(x), \quad \text{and} \quad (4)$$

$$H(g \circ f)(x) = g'(f(x))H + g''(f(x))\nabla f(x)\nabla f(x)^T \quad (5)$$

Note that here $d = \nabla f(x)$ is viewed as a column vector and so $\nabla f(x)\nabla f(x)^T$ is the *outer product*

$$\nabla f(x)\nabla f(x)^T = dd^T = (d_i d_j)_{ij}.$$

We will need the following example to construct test functions for unconstrained minimization:

Example 2.1. Let $\phi : G \subseteq \mathbb{R} \rightarrow \mathbb{R}$ be a function of one variable defined on an open subset $G \subseteq \mathbb{R}$, fix $a \in \mathbb{R}^n$ and set $g(x) = \phi(a \cdot x)$, for all $x \in \mathbb{R}^n$ such that $a \cdot x \in G$.

Since $f(x) = a \cdot x$ satisfies $\nabla f(x) = a$ and $H = \nabla^2 f(x) = 0$, for all $x \in \mathbb{R}^n$, the formulas (4) and (5) yield

$$\nabla g(x) = \phi'(a \cdot x)a \quad \text{and} \quad \nabla^2 g(x) = \phi''(a \cdot x)aa'.$$

The idea is to construct test functions of the form

$$\text{obj}F(x) = \sum_j \phi_j(a_j \cdot x)$$

where all the $\phi_j = \phi_j(u)$ have a unique minimum at $u = 0$. Then $\text{obj}F(x)$ assumes its global minimum at all points x satisfying $a_j \cdot x = 0$, for all j , equivalently $Ax = 0$, where A is the matrix with rows a_j , provided such a solution exists.

If the functions ϕ_j are all convex, the same is true of our objective function $\text{obj}F$ (sums and compositions of convex functions are again convex) With this we can construct examples with well conditioned, poorly conditioned and even singular Hessians ($\ker(A) \neq \{0\}$) with known minimizers. The conditioning of $\nabla^2 f(x)$ is closely related to that of the matrix A .

3 Nullspace

Here we discuss methods to solve an underdetermined set of linear equations $Ax = b$ where the number n of variables is larger than the number m of equations. In other words A is an $m \times n$ matrix with $m < n$.

The set of solutions is then the hyperplane $x_0 + \ker(A)$ where $\ker(A)$ denotes the nullspace of A and x_0 is any particular solution of $Ax = b$.

We will represent the nullspace of A as the column space $colspace(F)$ of a matrix F such that

$$AF = 0.$$

If F satisfies this equation, then the column space of F is a subspace of the null space $ker(A)$. Recall also that the column space of F is the range $Im(F)$ of the linear map defined by the matrix F .

We will generally assume that A has full rank, i.e. $rank(A) = m$. The nullspace $ker(A)$ then has dimension $n - m$ and the columns of F span the entire nullspace of A if and only if $rank(F) = n - m$.

If we have found such a matrix F and x_0 with $Ax_0 = b$, then the hyperplane of all solutions to $Ax = b$ can be represented as

$$x_0 + ker(A) = x_0 + Im(F).$$

We want to apply this to the minimization problem

$$\text{minimize } f : \mathbb{R}^n \rightarrow \mathbb{R} \text{ under the constraint } Ax = b. \quad (6)$$

Since the solutions x to the constraint are exactly the vectors x of the form $x = x_0 + Fu$, $u \in \mathbb{R}^{n-m}$, we can turn this into the unconstrained problem

$$\text{minimize } g(u) = f(x_0 + Fu) \text{ on all of } \mathbb{R}^{n-m}. \quad (7)$$

Thereby we reduce the dimension of the problem from n to $n - m$. In large scale problems one would not do this since this operation can destroy the sparseness of the Hessian $H(f)$ which is critical for computation in very large dimensional problems.

In dense small problem (up to say $n = 3000$ variables) the above elimination of the constraint $Ax = b$ is a reasonable approach. We will discuss two algorithms to finding F and x_0 . In both cases the columns of F will be orthonormal and hence an orthonormal basis for the null space $ker(A)$.

3.1 Null space with QR-decomposition of A

This approach relies on the relation $ker(A) = Im(A')^\perp$, where the prime denotes the matrix transpose and V^\perp is the orthogonal complement of a subspace V . Recall that

$$\begin{aligned} n &= dim(ker(A)) + dim(ker(A)^\perp) = dim(ker(A)) + dim(Im(A')) \\ &= (n - m) + dim(Im(A')). \end{aligned}$$

Thus $dim(Im(A')) = m$. We will find the range $Im(A')$ and its orthogonal complement using the QR-decomposition of the transpose A' (which has dimension $n \times m$, i.e. maps from \mathbb{R}^m to \mathbb{R}^n):

$$A' = QR, \quad (8)$$

where Q is orthogonal and R upper triangular with nonzero diagonal, where R is $p \times m$ (i.e. maps from $\mathbb{R}^m \rightarrow \mathbb{R}^p$ and Q is $m \times p$, i.e maps from $\mathbb{R}^p \rightarrow \mathbb{R}^m$).

In such a factorization Q and R must be $n \times p$ and $p \times m$ respectively, since A' is $n \times m$. Moreover $\text{rank}(Q) \geq \text{rank}(A') = \text{rank}(A) = m$ and so we must have $p \geq m$. Because the columns of Q are linearly independent and of dimension n we also must have $p \leq n$.

Indeed such factorizations exist for all $m \leq p \leq n$ but in practice (i.e. available in libraries) there are only two flavours:

(A) The reduced (minimal) form gives us R as an $m \times m$ matrix and $Q = [q_1, \dots, q_m]$ as an $n \times m$ matrix with m -columns $q_j \in \mathbb{R}^n$. It follows that R maps onto $\mathbb{R}^m = \text{dom}(Q)$ and hence

$$\text{Im}(A') = \text{Im}(Q) = \text{span}(q_1, \dots, q_m).$$

From this we see that $\ker(A) = \text{Im}(A')^\perp = \text{span}(q_1, \dots, q_m)^\perp$ but if we use the reduced form we have to compute this orthogonal complement ourselves. In other words we have to extend $\{q_1, \dots, q_m\}$ to an orthonormal basis $\{q_1, \dots, q_m, q_{m+1}, \dots, q_n\}$ of \mathbb{R}^n ourselves.

In R the minimal form is obtained as follows

```
qrA <- qr(t(A)); Q <- qr.Q(qrA); R <- qr.R(qrA)
```

Here we first compute a QR-decomposition object `qrA` and from this object extract the matrices Q and R using the helper functions `qr.Q` and `qr.R`.

The intermediate forms of the decomposition (8) with $m \leq p \leq n$ extend this matrix Q by adding (arbitrarily) orthonormal columns on the right

$$Q \rightarrow Q_+ = [q_1, \dots, q_m, q_{m+1}, \dots, q_p]$$

and adjusting the matrix R by adding zero rows at the bottom (so that R becomes $p \times m$). This of course does not change the matrix product QR as we can see from block multiplication

$$Q_+ = [Q, Q_1], \quad R_+ = \begin{pmatrix} R \\ 0 \end{pmatrix} \quad \implies \quad Q_+ R_+ = QR + Q_1 * 0 = QR.$$

(B) The full (maximal) form of the decomposition has $p = n$, that is, the columns of Q have been extended to a full orthonormal basis of \mathbb{R}^n and so clearly

$$\ker(A) = \text{Im}(A')^\perp = \text{span}(q_1, \dots, q_m)^\perp = \text{span}(q_{m+1}, \dots, q_n). \quad (9)$$

In other words the matrix F can be chosen to be the matrix with columns q_{m+1}, \dots, q_n :

$$F = [q_{m+1}, \dots, q_n] \in \text{Mat}_{n \times (n-m)}(\mathbb{R}). \quad (10)$$

To get a special solution x_0 of $Ax = b$ note that $A = R'Q'$ and solve the factored form $R'Q'x = b$. Set $y = Q'x$. Then *forward solve* the lower triangular system $R'y = b$ and get x from $Q'x = y$ as $x = Qy$.

To get the complete factorization $A' = Q_+ R_+$ in R we do

```
qrA <- qr(t(A)); Qp <- qr.Q(qrA,complete=TRUE); R <- qr.R(qrA)
```

and then extract the matrix Q via $Q <- Qp[,1:m]$. R is already in incomplete form (no zero rows at bottom) since we have not specified `complete=TRUE` in the function `qr.R` (the default is `complete=FALSE`).

With this the factorization of A' becomes $A' = QR$ (and not $A' = QpR$) and this is what we need in the computation of the special solution x_0 above.

The same holds true in the scala `breeze` library where we get the `QR` decomposition as

```
val qr.QR(q,r)=qr(A).
```

This yields $Q = q$ as $n \times n$ -matrix but yields $R = r$ as $m \times m$ -matrix, that is, without the $n - m$ bottom zero rows. With this the decomposition of A' becomes $A' = Q[\cdot, 1 : m] * R$ and so $A = R'P'$ where the matrix

$$P = Q[\cdot, 1 : m] = [col_1(Q), \dots, col_m(Q)]$$

has orthogonal columns and we must solve for x_0 as $R'y = b$, $P'x = y$, that is, $x = Py$. Let us note that this x_0 is the minimal norm solution of the system $Ax = b$. Indeed $x_0 = Py \in span(col_1(Q), \dots, col_m(Q))$ while $Fu \in span(col_{m+1}(Q), \dots, col_n(Q))$ and so $Fu \perp x_0$, for all $u \in \mathbb{R}^{n-m}$, from which the claim follows.

The `QR` factorization is computed by repeatedly applying (orthogonal) Householder updates and is thus a very stable algorithm.

3.2 Nullspace via SVD decomposition of A

We can also compute the nullspace $ker(A)$ and special solution x_0 of $Ax = b$ using the more involved *SVD-decomposition* of A :

$$A = U\Sigma V'$$

where $U \in Mat_{m \times m}(\mathbb{R})$ and $V \in Mat_{n \times m}(\mathbb{R})$ are orthonormal matrices and Σ is an $m \times n$ diagonal matrix with entries

$$\sigma_1 \geq \sigma_2 \geq \dots \geq \sigma_m \geq 0$$

(the singular values of A). Since $rank(A) = m$ we must have $\sigma_i > 0$ and so the matrix Σ is invertible. Write $V = [v_1, \dots, v_m]$, where the $v_j \in \mathbb{R}^n$ are the columns of V . Thus, for all $x \in \mathbb{R}^n$,

$$\begin{aligned} Ax = 0 & \iff U\Sigma V'x = 0 \\ & \iff \Sigma V'x = 0 \\ & \iff 0 = V'x = (v_1 \cdot x, v_2 \cdot x, \dots, v_m \cdot x)' = 0 \\ & \iff x \perp \{v_1, \dots, v_m\} \\ & \iff x \in span(v_{m+1}, \dots, v_n) \end{aligned}$$

In other words

$$\ker(A) = \text{span}(v_{m+1}, \dots, v_n). \quad (11)$$

where $v_{m+1}, \dots, v_n \in \mathbb{R}^n$ are vectors which extend v_1, \dots, v_m to an ON-basis of \mathbb{R}^n . To get the full $n \times n$ matrix

$$V_+ = [V, v_{m+1}, \dots, v_n]$$

we must compute the SVD in R via

$$\text{svdA} \leftarrow \text{svd}(A, \text{nv}=\text{n}); \text{Vp} \leftarrow \text{svdA}\$v; V \leftarrow V[, 1:\text{m}]$$

Now a particular solution of $U\Sigma V'x = Ax = b$ can be found by solving $\Sigma V'x = U'b := c$ which is equivalent to

$$V'x = (c_1/\sigma_1, \dots, c_m/\sigma_m)',$$

A particular solution of this is given by

$$x_0 = Vw, \quad \text{where } w = (c_1/\sigma_1, \dots, c_m/\sigma_m)' \in \mathbb{R}^m. \quad (12)$$

since the columns of V are orthonormal.